



Use AI agents to improve your business

10 Jun 2026



- Executive summary
 - The business question, answer and implications
 - How can I use AI to improve my business, be more productive — and make more money?
- Business impact, Financial and Strategic View
- Questions we answer
 - Can AI make my business better? — Hours/days, not weeks/months
 - Which agents enforce quality rules? — the verdict
 - Which agent/model should lead?
 - What should leaders do next? — Turn the experiment into a capability

- This presentation and project deliverables was created with the assistance of Artificial Intelligence. All business scenarios, datasets, analyses, and visualizations are based on synthetic data and code generated by AI for demonstration purposes only.
- Its purpose is to showcase how modern AI technologies are becoming strategic business enablers, helping organizations make better decisions, operate more efficiently, uncover new opportunities, and deliver greater value to customers and stakeholders.
- All code of this project are available at GitHub
- **## Benchmark and Reproducibility Notes**
- ⚠ This is a small, controlled benchmark. LLMs are probabilistic — prompt, context, and evaluation changes can shift the ranking. Treat these results as a point-in-time comparison, not a definitive ranking. Expert review remains essential before any AI-generated solution goes to production.
- Each Agent/LLM combination included in this repository provides its own implementation artifacts, documentation, and generated outputs. These files are preserved to demonstrate the specific solution produced by that particular setup.
- When executing the project using the provided `instructions.md` file, your results may differ from the examples included in this repository. Variations can occur due to differences in agent harnesses, MCP server releases, skill definitions, or execution environments. Nevertheless, because all implementations leverage the same core skills and MCP server capabilities, they consistently converge on a similar architectural pattern.

THE NEXT SLIDES WERE GENERATED USING AI

AI - Build a solution
from scratch

From Sales Transactions to **Insights**

How a skilled person, paired with an AI coding agent, compresses weeks of consulting work into hours / days — for with \$3 maximum of compute per model.



A medallion data pipeline, built & deployed by AI on Databricks

“

How can I use AI to improve my business, be more productive — and make more money?

— *the question every CEO is asking*



The starting answer: put the knowledge your team already owns to work with AI.

AI doesn't replace your experts — it amplifies them, turning their know-how into working solutions far faster than before.

WHY NOW

The economics of a proof-of-concept just changed



THE TRADITIONAL POC

Discover business rules in workshops

- Weeks of meetings to define business rules
- Manual environment setup & configuration
- Hand-built pipelines, hand-written specs
- Specialist teams, sequential hand-offs

TIMELINE **Weeks → Months**



THE AI-ACCELERATED POC

One person + one AI coding agent


- Agent reads the data & infers the rules
- Designs schema, pipeline & quality checks
- Writes the code and deploys it end-to-end
- Generates the specs & docs as it builds

TIMELINE **Hours → Days**

THE EXPERIMENT

A real consulting example: raw data to insight

The same single instruction file and the same task were given independently to seven frontier AI coding agents. Each one read a 100,000-row sales dataset, inferred the business rules, then designed and deployed a Bronze → Silver → Gold pipeline on Databricks — end to end, with no human guidance during the build.



THE LINE-UP

Seven frontier AI agents, one identical brief


Coding Agent	Frontier LLM	Deployment Method
GitHub Copilot	Claude Sonnet 4.6	Direct build in workspace
Claude Code	Claude Opus 4.8 ★	Direct build in workspace
Claude Code	Claude Fable 5	Direct build in workspace
GitHub Copilot	GPT-5.3 Codex	Databricks Asset Bundle (DAB)
OpenAI Codex	GPT-5.5	Databricks Asset Bundle (DAB)
GitHub Copilot	GPT-5.4	Databricks Asset Bundle (DAB)
Databricks Genie	Databricks-native (Genie)	Direct build in workspace

★ *Claude Opus 4.8 was ultimately selected as the foundation — see the verdict.*

KEY FINDING #1

Where the models diverged: data quality

From the same 100,000 raw rows, five agents enforced data-quality rules — invalid ages, negative prices, nulls, impossible ship dates — while two passed every row through unfiltered.



Rows retained in the Silver layer (of 100,000 raw)

Action item: don't assume a model will look for data-quality issues on its own — especially lower-cost models. Make remediation an explicit instruction.

KEY FINDING

The verdict — and the caveat



Claude Opus 4.8 — selected

- Strongest data-quality refinement overall
- Best documentation & final deliverables
- Strong price-to-quality vs. Fable 5
- Cleansed data while preserving more good rows
- Now the “Agent Manager” coordinating audit & auto-ingestion



Read this before you act

This is a small, controlled benchmark. LLMs are probabilistic — a better prompt could put GPT-5.5 on top instead. Treat it as a point-in-time comparison, not a definitive ranking. **Expert review remains essential before anything ships to production.**

THE BOTTOM LINE

Why this matters to the business



Speed

An entire data-platform PoC — schema, pipeline, quality logic, deploy, validation — scaffolded in hours, not weeks.



Cost

About \$3 of compute per run to produce a working, deployed solution worth weeks of specialist effort.



Leverage

One skilled person steers the agent through discovery, design, code and deployment — output of a whole team.



Guardrail

None of the outputs ship as-is. AI accelerates the work; your experts review, combine and harden it.

A single instruction file plus a frontier-model agent can scaffold a data platform from scratch — a real accelerator for data & analytics teams.

From experiment to operating capability

1



Executive dashboards

CEO/CFO metrics and tables on the curated Gold layer — including a simple forecast tab.

2



Automated audit

Generated SQL (and Genie at the workspace) tallies every table & view the pipeline created.

3



Autonomous ingestion

New data flows in without manual setup, with Opus 4.8 acting as the Agent Manager.

Your team's knowledge + AI agents = a data intelligence platform in days, not quarters.

END OF AI-GENERATED SLIDES

PROJECT DELIVERABLES

Databricks Dashboards and Insights examples



Sales Performance — Executive Dashboard (cowork_op48) ★

[Edit draft](#)[21h ago](#) [Schedule](#) [Share](#)[Executive Overview](#) [Operations & Geography](#) [Forecast](#)

Sales Performance — Executive Overview

Revenue, order, and customer KPIs across the cowork_op48 gold layer (Jan 2023 – present)

Total Net Revenue

\$ 25,02 mi

Total Orders

11,33 mil

Avg Order Value


2,21 mil

Distinct Customers


11 mil

Revenue Trend & Category Mix

Monthly Net Revenue Trend



Revenue Share by Category



Databricks Dashboards and Insights examples



Sales Performance — Executive Dashboard (cowork_op48) ★ [Edit draft](#)

⋮ 21h ago [Schedule](#) [Share](#)

[Executive Overview](#) [Operations & Geography](#) [Forecast](#)

Operations & Geography

Profitability drivers, fulfillment health (returns/cancellations), and regional performance

Total Gross Revenue

\$ 31,51 mi

Avg Discount Rate

20,7%

Return Rate


33,9%

Cancellation Rate

32,5%

Regional & Fulfillment Mix

Net Revenue by State



Order Mix by Status



Net Revenue by Payment Type

[Ask Genie](#)

Databricks Dashboards and Insights examples



Sales Performance — Executive Dashboard (cowork_op48) ★

[Edit draft](#)[21h ago](#) [Schedule](#) [Share](#)[Executive Overview](#) [Operations & Geography](#) [Forecast](#)

Revenue Forecast

Simple linear-trend forecast of monthly net revenue, fitted on 36 months of complete history (Jan 2023 – Dec 2025) and projected 6 months forward

Forecasted Revenue (Next 6 Months)

\$ 4,19 mi

Forecasted Avg Monthly Revenue


\$ 697,76 mil

Trend Change per Month

506,17

Actual vs. Trend / Forecast

Net Revenue: Actual vs. Trend / Forecast (6-month projection)

[Ask Genie](#)

Forecast Detail by Month



Databricks – data pipeline deployed using AI




Jobs & Pipelines >

pipeline_sales_medallion_fa5_v1 ☆ [Send feedback](#)

Run: Jun 09, 2026, 08:33 PM • Completed [Select tables for refresh](#) ⓘ

[Edit pipeline](#) [Settings](#) [Schedule](#) [Share](#) [Run pipeline](#) ▼



Tables 14 Performance 9

[Type](#) ▼ [Status](#) ▼

Name ↕	Catalog	Schema	Type	Duration	Output rec...	Expectations	Dropped	Warnings	Failed	
fact_sales	cowork_fable5	gold_fa5_v1	Streaming table	15s	-	-	-	-	-	
silver_sales_transactions	cowork_fable5	silver_fa5_v1	Streaming table	4s	11K	6 met 3 unmet	89K	2.2K	-	
dim_customer_src	N/A	N/A	View	-	-	-	-	-	-	
dim_location_src	N/A	N/A	View	-	-	-	-	-	-	




Current frontier models benchmarks – Additional info

JUN 2026

Highlights


Intelligence

Artificial Analysis Intelligence Index - Higher is better



Speed


Output tokens per second - Higher is better



Price


Updated

USD per 1M tokens (blended) - Lower is better



Artificial Analysis

- Intelligence
- Artificial Analysis Coding Agent Index
- Image & Video Leaderboards
- Intelligence Breakdown
- AA-Omniscience
- GDPval-AA
- ITBench-AA · NEW
- Artificial Analysis Openness Index
- Output Tokens
- Cost Efficiency
- Speed & Latency
- Price · Updated
- API Provider Performance



- Intelligence
- Artificial Analysis Coding Agent Index
- Image & Video Leaderboards
- Intelligence Breakdown
- AA-Omniscience
- GDPval-AA
- ITBench-AA · NEW
- Artificial Analysis Openness Index
- Output Tokens
- Cost Efficiency
- Speed & Latency
- Price · Updated
- API Provider Performance

Artificial Analysis Intelligence Index

Coding Index

Agentic Index


Artificial Analysis Intelligence Index

Artificial Analysis Intelligence Index v4.0 incorporates 10 evaluations: GDPval-AA, τ^2 -Bench Telecom, Terminal-Bench Hard, SciCode, AA-LCR, AA-Omniscience, IFBench, Humanity's Last Exam, GPQA Diamond, CritPt

28 of 536 models

+ Add model from specific provider

Artificial Analysis



- Intelligence
- Artificial Analysis Coding Agent Index
- Image & Video Leaderboards
- Intelligence Breakdown
- AA-Omniscience
- GDPval-AA
- ITBench-AA · NEW
- Artificial Analysis Openness Index
- Output Tokens
- Cost Efficiency
- Speed & Latency
- Price · Updated
- API Provider Performance

Artificial Analysis Intelligence Index Coding Index Agentic Index

Artificial Analysis Agentic Index


Represents the average of agentic capabilities benchmarks in the Artificial Analysis Intelligence Index (GDPval-AA, τ^2 -Bench Telecom)

28 of 536 models

+ Add model from specific provider



Artificial Analysis



Artificial Analysis

Intelligence vs. Cost to Run Artificial Analysis Intelligence Index

[Models](#)[Coding Agents](#)[Speech, Image, Video, Music](#)[Hardware](#)[Leaderboards](#)[About](#)[AI Trends](#)[Arenas](#)[Log In](#)[Q](#) [K](#)[X](#)[in](#)[YouTube](#)

28 of 536 models

[Add model from specific provider](#)

Most attractive quadrant

● OpenAI ● Google ● Anthropic ● Mistral ● DeepSeek ● xAI ● Amazon ● MiniMax ● NVIDIA ● Kimi ● Xiaomi ● Alibaba

